

UC Berkeley Data8

JupyterHub 2016 mini-workshop

Background

- Foundations of Data Science (Stat/CS/Info 8)
 - has “connector” courses
 - for students with no knowledge of Statistics or Computer Science
- Fall 2015: ~100 students; Hub was on bare metal instance
- Spring 2016: ~480 students; Hub was on Azure with fallback to previous instance
- Fall 2016: ~480 students
- Future: 500-1500?

Deployment

- Copy Jess' compmodels deployment
 - <https://github.com/compmodels/jupyterhub-deploy>
 - Currently without nbgrader
- Tested on AWS; took a couple of months to learn about all of the components.
- Learning curve, large ecosystem.
 - ansible, docker, swarm, dockerspawner, swarmspawner, docker-py, jupyter, notebooks, node, nginx, web sockets.

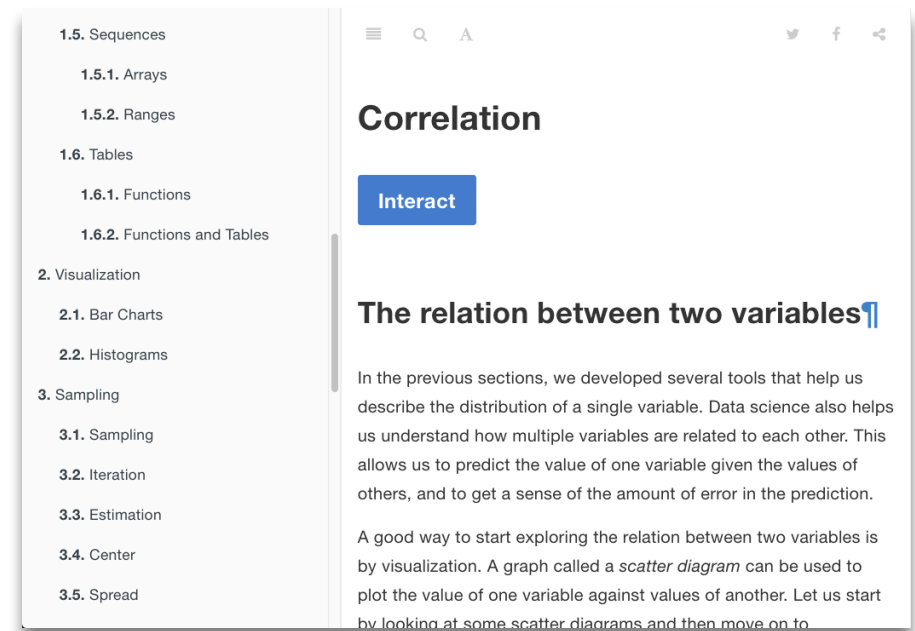
Versioning

- Versioning, dependencies, reproducibility...
 - A build process may break depending on when you execute it. Tools and libraries are always changing.
 - Be cautious about untagged “docker pull”, “pip install”. Wherever possible, embed git tags, docker image tags, conda and pip package versions, etc.
- Continuous Integration
- Reference deployments
 - <https://github.com/jupyterhub/jupyterhub-deploy-teaching>
 - <https://github.com/jupyterhub/jupyterhub-deploy-docker>

Content Distribution

How can faculty or graduate student instructors push notebooks or other files into student accounts?

- Interact: hub side service that pulls notebooks from github repos when students click on a textbook link.
- `/hub/interact?repo=foo&path=a/b.ipynb`
- Not all instructors are familiar with git.
- How to merge faculty changes with students' changes. Cell-centric contents manager? (Postgres)
- jupyter-drive: we can take advantage of UCB's Google Apps integration.



<http://www.inferentialthinking.com>

Content Development

- Instructors want to add libraries to our system user container.
- How to test, given that users aren't root in the container?
- We have a derived system user container meant to be run on laptops with Docker Beta
- imagespawner: HTML form that allows the user to choose which container they want to run. Give instructors a choice between regular container and one where the user has elevated privileges
 - <https://github.com/ryanlovet/containerspawner>

Scaling

- How to support multiple courses
 - Data8 has “connectors”
- One hub with many courses or one course per hub instance?
- How to scale compute?
 - Spawners can do it: Swarm, kubernetes
 - Cloud APIs: AWS, Azure. Have spawner provision additional nodes on demand.
- How well does JupyterHub scale?
 - Start collecting data: statsd

Future Direction

- Keep going, knock on wood.
- Is our use case too large for JupyterHub? Should we use “NBaaS”?
- <http://notebooks.azure.com>
 - No need to create the infrastructure
 - Less control of server side
 - Requires each student have a cloud account
- <http://mybinder.org>

Links

- <http://data8.org>
- <http://data.berkeley.edu>
- <http://inferentialthinking.com>